

20-numbers-headaches

June 2, 2024

1 Numbers headaches

C++ benefits for scientific computing: * performance and precision is the primary priority * compatibility with C

Drawbacks: * portability is not the primary priority * some C inherited features, such as implicit conversions, may lead to tricky numerical errors.

1.1 Uninitialized variables

First thing to be very careful about: if a basic variable is defined without any initial value, there is no guarantee that its value is set to 0. The C designers probably thought it was a waste of precious runtime...

Recommendation : **always give an initial value to your variables.**

If by mistake you use an uninitialized variable, no doubt that the compiler will give you a warning.

Recommendation : **compiler's warnings should be taken with the utmost seriousness.**

1.2 Unportable numeric types

The size of numeric variable types in C++ depends on the implementation. This may impede the portability of the code.

For example, the following rules are imposed on integer types by the C++ standards: * **short** : a width of at least 16 bits. * **int** : a width of at least 16 bits. * **long** : a width of at least 32 bits. * `sizeof(short) <= sizeof(int) <= sizeof(long)`

The rules on floating point types are not strict either: * `sizeof(float) <= sizeof(double) <= sizeof(long double)` * **float** : typically 32 bits (*IEEE 754, 6-9 significant digits, typically 7*). * **double** : typically 64 bits (*IEEE 754, 15-18 significant digits, typically 16*). * **long double** : 80 to 128 bits (*18-36 significant digits*).

For a given platform, the standard class `numeric_limits` can help to check the sizes:

```
[6]: #include <iostream>
```

```
[7]: #include <limits>
```

```
std::cout  
<< "type\tbits\tmin\t\tmax\n"
```

```
<< "int\t" << sizeof(int)*8 << "\t"
<< std::numeric_limits<int>::min() << '\t'
<< std::numeric_limits<int>::max() << '\n'
<< "float\t" << sizeof(float)*8 << "\t"
<< std::numeric_limits<float>::min() << '\t'
<< std::numeric_limits<float>::max() << '\n'
<< "double\t" << sizeof(double)*8 << "\t"
<< std::numeric_limits<double>::min() << '\t'
<< std::numeric_limits<double>::max() << '\n' ;
```

type	bits	min	max
int	32	-2147483648	2147483647
float	32	1.17549e-38	3.40282e+38
double	64	2.22507e-308	1.79769e+308

1.3 New in C++11: fixed size integer types

C++11 borrows from C99 a bunch of integer types following different rules: - good for memory, the smallest with at least N bits: `int_least8_t`, `int_least16_t`, `int_least32_t`, `int_least64_t` ; - good for performance, the fastest with at least N bits: `int_fast8_t`, `int_fast16_t`, `int_fast32_t`, `int_fast64_t` ; - good for portability, with exactly N bits: `int8_t`, `int16_t`, `int32_t`, `int64_t`.

They are available when including `<cstdint>`.

1.4 Hardly predictable precision

When an operation mix variables with different precisions, the compiler casts all the operands to the *best* precision.

```
[6]: float f = std::numeric_limits<float>::max() ;
std::cout<<( f * 10.f )<<std::endl ;
std::cout<<( f * 10. )<<std::endl ;
```

```
inf
3.40282e+39
```

Floating point types are considered better than integer types.

```
[10]: float f = std::numeric_limits<float>::max() ;
std::cout<<( f * 10 )<<std::endl ;
```

```
inf
```

unsigned types are considered better than their signed flavor...

```
[11]: unsigned int i = 1 ;
int j = -1 ;
std::cout<<( i * j )<<std::endl ;
```

```
4294967295
```

Sometimes, computation is performed with a precision higher than the original variables. For example, a `short` based operation will always be evaluated as an `int` (because `int` is the same size as the hardware registers).

```
[12]: short s1 = std::numeric_limits<short>::max() ;
      short s2 = 1 ;
      std::cout << (s1+s2) << std::endl ;
```

32768

```
[13]: short s3 = (s1+s2) ;
      std::cout << s3 << std::endl ;
```

-32768

BEWARE: Intel processors were typically computing their double operations with extraneous digits (80). One may think **the higher the precision, the better**. But it implies a **portability issue**, because the results will differ when running your code on a different processor, when vectorizing, when porting the code to GPU...

1.5 Unnoticed conversions

For ease of writing code, C compilers and thus, C++ compilers as well, are authorized to perform multiple conversions between predefined numeric types. These so-called **implicit conversions** are automatic and often unnoticed by the developer.

Some of these allowed implicit conversions can introduce a loss of precision: for example a transformation from a floating-point number to an integer. The compiler assumes that such **narrowing** is done on purpose. Is it still a reasonable assumption for a code made of thousands of lines?

```
[15]: double pi = 3.1416 ;
      int i = pi ;
      std::cout<<"double "<<pi<<" => int "<<i<<std::endl ;
```

double 3.1416 => int 3

```
[16]: long lmax = std::numeric_limits<long>::max() ;
      short s = lmax ;
      std::cout<<"long "<<lmax<<" => short "<<s<<std::endl ;
```

long 9223372036854775807 => short -1

```
[17]: double dmax = std::numeric_limits<double>::max() ;
      float f = dmax ;
      std::cout<<"double "<<dmax<<" => float "<<f<<std::endl ;
```

double 1.79769e+308 => float inf

Even worse, the compiler can transform any signed/unsigned integer into unsigned/signed !

```
[18]: void display_signed( short v )
      { std::cout<<v<<std::endl ; }
```

```
unsigned short us = 42000 ;
display_signed(us) ;
```

-23536

```
[19]: void display_unsigned( unsigned short v )
      { std::cout<<v<<std::endl ; }

      short s = -42 ;
      display_unsigned(s) ;
```

65494

BEWARE: when mixing signed and unsigned integers in an expression, the compiler will consider the unsigned flavor as the most accurate, and transform all the integers accordingly. One more time, **paying attention to compiler warnings will help you out.**

1.6 Disputed practice: never use unsigned numbers

...if you can ! - There are some contexts (embedded computing) where every bits is worth saving. - The standard library designers made the choice of unsigned integers for the size and indexes of all the containers :(

1.7 Good old-fashioned practice: make all conversions explicit

In a program of a large size, implicit conversions are more of a hindrance than a help. It is advised to: - set the warning level to maximum, - scrutinize carefully all compiler warnings, - make explicit any conversion you identify in the code.

The C way, for explicit conversions, is to use the type name as a function:

```
[17]: unsigned short i = 42000 ;
      short j = short(i) ;
      std::cout<<j<<std::endl ;
```

-23536

But this does not catch the eye. Better, C++ comes with a set of explicit type casting operators. The one to be used by default is `static_cast`:

```
[10]: unsigned short i = 42000 ;
      short j = static_cast<short>(i) ;
      std::cout<<j<<std::endl ;
```

-23536

Three other type casting operators are available, for rare specific use-cases: * `const_cast` : in rare cases, when one wants to get rid of the constness of a variable; * `dynamic_cast` : to goes down an inheritance tree; * `reinterpret_cast` : in very rare cases, when one wants to change the way a memory chunk is interpreted.

1.8 User-defined numerical types

If you are daredevil enough to develop your own numerical type, you may want to provide a constructor and/or a conversion operator, so to ease interaction with functions which generate and/or require doubles.

restart kernel

```
[13]: #include <iostream>
```

```
[14]: class Number
{
    public :
        Number( double f ) { std::cout<<"double_to_number"<<std::endl ; }
        operator double() { std::cout<<"number_to_double"<<std::endl ; return 0 ; }
} ;
```

```
[15]: Number number_pi(3.14) ;
```

double_to_number

```
[16]: void foo_double( double ) {}
```

```
[17]: foo_double(number_pi) ;
```

number_to_double

```
[18]: void foo_number( Number ) {}
```

```
[19]: foo_number(3.14) ;
```

double_to_number

BEWARE: the unary constructor, and the conversion operator, opens the door for **implicit conversions**. The compiler can even chain several ones, such as `short ==> double ==> Number` below, or the contrary.

```
[20]: short sh = 42 ;
      foo_number(sh) ;
```

double_to_number

```
[21]: void foo_short( short ) {}
```

```
[22]: display_short(number_pi) ;
```

number_to_double

1.9 Good old-fashioned practice: make unary constructors explicit

The implicit conversions problem does not only apply to numerical classes. It is true for any class which has a unary constructor (constructor with only one argument).

The keyword `explicit` forbids the use of those unary constructors for implicit conversions. It should be used almost always.

restart kernel

```
[1]: #include <iostream>
```

```
[2]: class Number
    {
    public :
        explicit Number( double f ) { std::cout<<"double_to_number"<<std::endl ; }
        operator double() { std::cout<<"number_to_double"<<std::endl ; return 0 ; }
    } ;
```

```
[3]: void foo_number( Number ) {}
```

```
[4]: foo_number(42) ;
```

```
input_line_10:2:2: error: no matching function for call
```

```
to 'foo_number'
```

```
foo_number(42) ;
```

```
^~~~~~
```

```
input_line_9:1:6: note: candidate function not viable:
```

```
no known conversion from 'int' to '__clang_N52::Number' for 1st argument
```

```
void foo_number( Number ) {}
```

```
^
```

Interpreter Error:

Yet the implicit use of the conversion operator is still ok.

```
[5]: Number number_pi(3.14) ;
```

```
double_to_number
```

```
[6]: void foo_double( double ) {}
```

```
[7]: foo_double(number_pi) ;
```

```
number_to_double
```

1.10 New in C++11 : explicit conversion operators

The keyword `explicit` has been generalized to conversion operators within C++11.

restart kernel

```
[1]: #include <iostream>
```

```
[2]: class Number
    {
    public :
        explicit Number( double f ) { std::cout<<"double_to_number"<<std::endl ; }
        explicit operator double() { std::cout<<"number_to_double"<<std::endl ;
        ↪return 0 ; }
    } ;
```

```
[3]: void display_double( double ) {}
```

```
[4]: Number value(3.14) ;
```

double_to_number

```
[5]: display_double(value) ;
```

input_line_11:2:2: **error:** no matching function for call

to 'display_double'

display_double(value) ;

~~~~~

input\_line\_9:1:6: note: candidate function not viable:

no known conversion from '::\_\_clang\_N52::Number' to 'double' for 1st argument

void display\_double( double ) {}

^

Interpreter Error:

## 2 Take Away

- I use *unsigned integers* only for sizes and indexes.
- Implicit conversions are a major source of bugs in ancient C++.
- Modern C++ provides different ways to take the control back:
  - explicit constructors and converters,
  - forbidden overloads (to be seen later),
  - **universal initialization**.

## 3 Questions ?

## 4 References

Types \* [Cpp Reference](#) \* [Fixed size integer types](#)

Implicit conversions - [Cpp Reference](#) - [LearnCPP](#) - [Nothing But](#)

© CNRS 2024

*This document was created by David Chamont and translated by Olga Abramkina. It is available under the [Licence Creative Commons - Attribution - No commercial use - Shared under the conditions 4.0 International](#)*