

4-execution

June 2, 2024

1 New execution strategies in C++17/20

C++17 takes a further step in the direction of automatic code parallelization by adding a new argument, the **execution policy**, to most of the algorithms in the standard library.

The available policies are `std::execution::seq`, `std::execution::par` and `std::execution::par_unseq`. C++20 is adding `std::execution::unseq`.

WARNING: the exceptions are not compatible with the algorithms which takes an execution policy as first argument.

WARNING: the actual application of parallelization is highly dependent on the hardware and the implementation of the compiler and its standard library. This requires most of the time (if not always) a **backend technology**. For example, GCC relies on Intel TBB.

Let's browse a toy code which takes a collection of floating point numbers, raises them to a given power, then prints their mean value.

```
[1]: %%file tmp.execution-common.h

#include <vector>
#include <cmath>
#include <algorithm>
#include <numeric>
#include <iostream>
#include <cassert>

using Real = double ;
using Reals = std::vector<Real> ;
```

Overwriting tmp.execution-common.h

```
[2]: %%file tmp.execution-prepare-input.h

// random numbers in [1.-1./scal,1.]
void generate( Reals & rs, Real scale )
{
    srand(1) ;
    for ( auto & r : rs )
        { r = 1.-rand()/scale/RAND_MAX ; }
```

```
}
```

Overwriting tmp.execution-prepare-input.h

```
[3]: %%file tmp.execution-process.h

// compute  $x^{\text{degree}}$  and store it into ys
void pow
( Reals const & xs, int degree, Reals & ys )
{
    std::transform(xs.begin(),xs.end(),ys.begin(),
        [degree]( Real x )
        {
            Real y {1.0} ;
            for ( int d=0 ; d<degree ; ++d )
                { y *= x ; }
            return y ;
        }) ;
}
```

Overwriting tmp.execution-process.h

```
[4]: %%file tmp.execution-post-process-output.h

// compute the mean
void postprocess( Reals const & rs )
{ std::cout<<"mean: "<<(std::accumulate(rs.begin(),rs.end(),Real{0.})/rs.
    size())<<"\n" ; }
```

Overwriting tmp.execution-post-process-output.h

```
[5]: %%file tmp.execution.cpp

#include "tmp.execution-common.h"
#include "tmp.execution-prepare-input.h"
#include "tmp.execution-process.h"
#include "tmp.execution-post-process-output.h"

// main program
int main ( int argc, char * argv[] )
{
    assert(argc==3) ;
    std::size_t dim {std::stoul(argv[1])} ;
    int degree = atoi(argv[2]) ;

    // prepare input
    Reals input(dim) ;
    generate(input,degree) ;
```

```

// compute ouput
Reals output(dim) ;
pow(input,degree,output) ;

// post-process
postprocess(output) ;
}

```

Overwriting tmp.execution.cpp

```

[6]: %%file tmp.execution.sh
echo

rm -f tmp.execution.exe \
&& g++ -std=c++17 -O2 tmp.execution.cpp -o tmp.execution.exe\
&& time ./tmp.execution.exe $*

echo

```

Overwriting tmp.execution.sh

```

[9]: !bash -l tmp.execution.sh 1024 100000

```

mean: 0.62705

```

real    0m0.099s
user    0m0.099s
sys     0m0.000s

```

1.1 Parallel execution policy

When using an algorithm from the standard library, we can **suggest** a multi-threaded execution using a simple additional argument:

```

[10]: %%file tmp.execution-process.h

#include <execution>

// compute xs^degree and store it into ys
void pow
( Reals const & xs, int degree, Reals & ys )
{
    std::transform(std::execution::par,xs.begin(),xs.end(),ys.begin(),[degree](
↪Real x )
    {
        Real y {1.0} ;
        for ( int d=0 ; d<degree ; ++d )

```

```

    { y *= x ; }
    return y ;
  } ) ;
}

```

Overwriting tmp.execution-process.h

Yet, for the time being, with GCC, do not forget to add the library `-ltbb`, because Intel TBB is the backend implementation which enables GCC to apply `std::execution::par`.

The simplicity of the written code as a drawback : additional installations and compilation options, at least today.

```

[13]: %%file tmp.execution.sh
echo

rm -f tmp.execution.exe \
&& g++ -std=c++17 -O2 -ltbb tmp.execution.cpp -o tmp.execution.exe\
&& time ./tmp.execution.exe $*

echo

```

Overwriting tmp.execution.sh

```

[15]: !bash -l tmp.execution.sh 1024 100000

```

mean: 0.62705

```

real    0m0.008s
user    0m0.105s
sys     0m0.000s

```

1.2 Unsequenced execution policy

It seems just as easy to suggest the use of vectorized SIMD instructions, with the `std::execution::unseq` execution policy (C++20 only)...

```

[16]: %%file tmp.execution-process.h

#include <execution>

// compute xs^degree and store it into ys
void pow
( Reals const & xs, int degree, Reals & ys )
{
    std::transform(std::execution::unseq,xs.begin(),xs.end(),ys.begin(),[degree](
↪Real x )
    {

```

```

    Real y = 1.0 ;
    for ( int d=0 ; d<degree ; ++d )
        { y *= x ; }
    return y ;
}) ;
}

```

Overwriting `tmp.execution-process.h`

... but the speedup is not there. An important point is that the use of an execution policy argument **allows some parallelism, but does not make it mandatory**. It is an invitation given to the compiler, that it will accept or not, depending on the code context and the backends available.

```
[17]: !bash -l tmp.execution.sh 1024 100000
```

mean: 0.62705

```

real    0m0.100s
user    0m0.096s
sys     0m0.004s

```

1.3 Availability of this features (for what concerns g++)

The effective implementation of the new parallel execution strategies requires a background implementation.

The last version of g++ (11.2), together with Intel TBB as backend, automatically apply multi-threading when one ask a `std::execution::par` policy.

On the contrary, currently, g++ does not care about `std::execution::unseq`, and rather relies on its auto-vectorization feature.

2 Questions ?

2.1 Resources

- [Rainer Grimm's blog](#)
- [Rainer Grimm's blog](#)
- [Bartek's coding blog](#)
- [Compiler support](#) (see "Standardization of Parallelism TS")
- [Intel parallel STL](#)
- [Intel TBB](#)

© CNRS 2024

This document was created by David Chamont and translated by Olga Abramkina. It is available under the [License Creative Commons - Attribution - No commercial use - Shared under the conditions 4.0 International](#)