

1-arrays

June 3, 2024

1 Choose your data structure

1.1 AoS (Array of Structs)

In the code example below, a “SAXPY” ($y = a \cdot x + y$) calculation is done on a collection of XY elements.

```
[12]: %%file tmp.xy.h

struct XY
{
    double x, y {0.} ;
    void saxpy( double a )
    { y = a*x + y ; }
} ;
```

Overwriting tmp.xy.h

```
[13]: %%file tmp.aos-functions.h

#include <cstdlib> // for rand

template< typename Itr >
void randomize_x( Itr begin, Itr end )
{
    for ( Itr itr = begin ; itr!=end ; ++itr )
        { itr->x = std::rand()/(RAND_MAX+1.)-0.5 ; }
}

template< typename Itr >
void saxpy( Itr begin, Itr end, double a )
{
    for ( Itr itr = begin ; itr!=end ; ++itr )
        { itr->saxpy(a) ; }
}

template< typename Itr >
double accumulate_y( Itr begin, Itr end )
{
```

```

double res {0.} ;
for ( ltr itr = begin ; itr!=end ; ++itr )
{ res += itr->y ; }
return res ;
}

```

Overwriting tmp.aos-functions.h

```

[20]: %%file tmp.aos.cpp

#include "tmp.xy.h"
#include "tmp.aos-functions.h"
#include <cassert> // for assert
#include <cstdlib> // for atoi
#include <iostream>

int main( int argc, char * argv[] )
{
    assert(argc==3) ;
    int size {atoi(argv[1])} ;
    int repeat {atoi(argv[2])} ;
    std::cout.precision(18) ;

    XY * collection {new XY[size]} ;
    auto begin {collection} ;
    auto end {begin+size} ;

    randomize_x(begin,end) ;
    while (repeat--)
        saxpy(begin,end,0.1) ;
    double res {accumulate_y(begin,end)/size} ;
    std::cout<<res<<std::endl ;

    delete [] collection ;
}

```

Overwriting tmp.aos.cpp

```

[18]: %%file tmp.aos.bash
echo

rm -f tmp.aos.exe tmp.aos.py
g++ -std=c++17 tmp.aos.cpp -o tmp.aos.exe
./tmp.aos.exe $*

echo "s = 0" >> tmp.aos.py
for i in 0 1 2 3 4 5 6 7 8 9
do \time -f "s += %U" -a -o ./tmp.aos.py ./tmp.aos.exe $* >> /dev/null

```

```
done
echo "print('~ {:.3f} s').format(s/10.)" >> tmp.aos.py
python3 tmp.aos.py

echo
```

Overwriting tmp.aos.sh

```
[19]: !bash -l tmp.aos.bash 1024 100000
```

```
67.5053500207703507
(~ 1.380 s)
```

The `main` function is currently using an old-fashioned C array, and the script does not set explicitly the GCC optimization option, which means it is using the default `-O0` (no compiler optimization).

You are asked to try this code, then investigate the alternative arrays `std::array`, `std::valarray`, `std::vector`, `std::list` and the alternative GCC compilation options `-O2` (usual optimisations) and `-O3` (aggressive optimizations, including automatic vectorization). Fill the results below, and try to explain the differences.

Array	Option	-O0	-O2	-O3
Classic C array		0.	0.	0.
<code>std::array</code>		0.	0.	0.
<code>std::valarray</code>		0.	0.	0.
<code>std::vector</code>		0.	0.	0.
<code>std::list</code>		0.	0.	0.

1.2 SoA (Struct of Arrays)

Now let's try another approach: instead of creating a structure that groups together `x` and `y` and making it into an array (as it is naturally done on an object-oriented approach), let's try to make a global structure that contains an array of `x` on one hand, and an array of `y` on the other hand.

This is what the code skeleton below offers, again using C arrays and default `-O0`. Again, try alternative collections and compilation options. Fill the results table and explain.

```
[7]: %%file tmp.soa.h

#include "tmp.xy.h"

class SoA
{
public :
    SoA( int size ) : m_size(size), m_xs(new double[size]), m_ys(new
↪double[size]) {}
    ~SoA() { delete [] m_xs ; delete [] m_ys ; }
```

```

int size() { return m_size ; }
XY operator()( int indice ) const
{ return { m_xs[indice], m_ys[indice] } ; }
auto & xs() { return m_xs ; }
auto & ys() { return m_ys ; }
void saxpy( double a )
{
    for ( int i=0 ; i<m_size ; ++i )
        m_ys[i] = a*m_xs[i] + m_ys[i] ;
}
private :
    int m_size ;
    double * m_xs ;
    double * m_ys ;
} ;

```

Overwriting tmp.soa.h

```

[8]: %%file tmp.soa-functions.h

#include "tmp.soa.h"
#include <cstdlib> // for rand

void randomize_x( SoA & collection )
{
    for ( int i=0 ; i<collection.size() ; ++i )
        { collection.xs()[i] = std::rand()/(RAND_MAX+1.0)-0.5 ; }
}

double accumulate_y( SoA & collection )
{
    double res {0.} ;
    for ( int i=0 ; i<collection.size() ; ++i )
        { res += collection.ys()[i] ; }
    return res ;
}

```

Writing tmp.soa-functions.h

```

[9]: %%file tmp.soa.cpp

#include "tmp.soa-functions.h"
#include <iostream>
#include <cassert> // for assert
#include <cstdlib> // for atoi

int main( int argc, char * argv[] )
{

```

```

assert(argc==3) ;
int size {atoi(argv[1])} ;
int repeat {atoi(argv[2])} ;

SoA collection(size) ;
randomize_x(collection) ;
while (repeat--
    collection.saxpy(0.1) ;
double res = accumulate_y(collection)/size ;

std::cout.precision(18) ;
std::cout<<res<<std::endl ;
}

```

Writing tmp.soa.cpp

```

[10]: %%file tmp.soa.bash
echo

rm -f tmp.soa.exe tmp.soa.py
g++ -std=c++17 tmp.soa.cpp -o tmp.soa.exe
./tmp.soa.exe $*

echo "s = 0" >> tmp.soa.py
for i in 0 1 2 3 4 5 6 7 8 9
do \time -f "s += %U" -a -o ./tmp.soa.py ./tmp.soa.exe $* >> /dev/null
done
echo "print('{:.3f} s').format(s/10.)" >> tmp.soa.py
python3 tmp.soa.py

echo

```

Writing tmp.soa.sh

```

[11]: !bash -l tmp.soa.bash 1024 100000

```

67.5053500207703507
(0.920 s)

To help in the analysis, [GodBolt](#) can be used, which allows to observe the dose of “inlining”, or to look for the presence of vectorial instructions in assembly, such as `addpd` (Add Packed Doubles) or `mulpsd` (Multiply Packed Double).

Array	Option	-O0	-O2	-O3
Classic C array		0.	0.	0.
std::array		0.	0.	0.
std::valarray		0.	0.	0.

Array	Option	-O0	-O2	-O3
std::vector		0.	0.	0.
std::list		0.	0.	0.

© CNRS 2021 Assembled and written in french by David Chamont, translated by Karim Hasnaoui, this work is made available according to the terms of the [Creative Commons License - Attribution - NonCommercial - ShareAlike 4.0 International](#)